


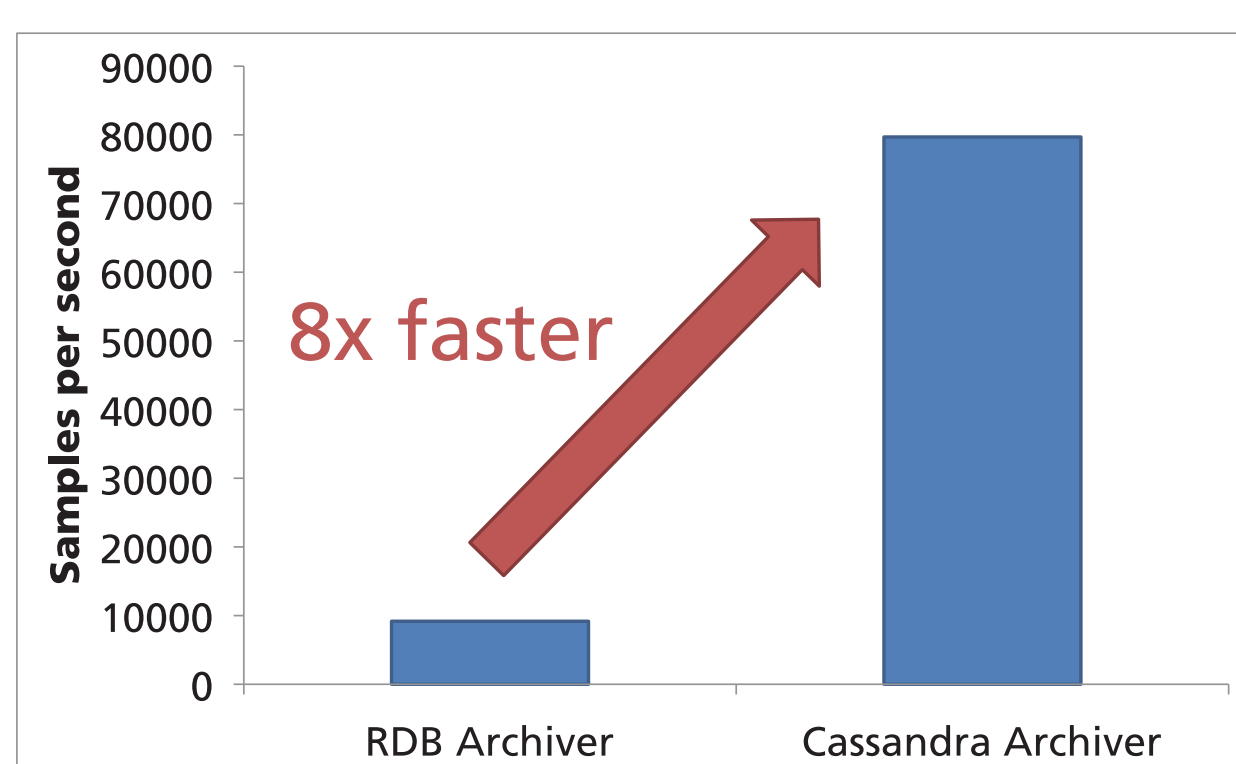
## Scalable Archiving with the Cassandra Archiver for CSS

S. Marsching, aquenos GmbH, Baden-Baden, Germany

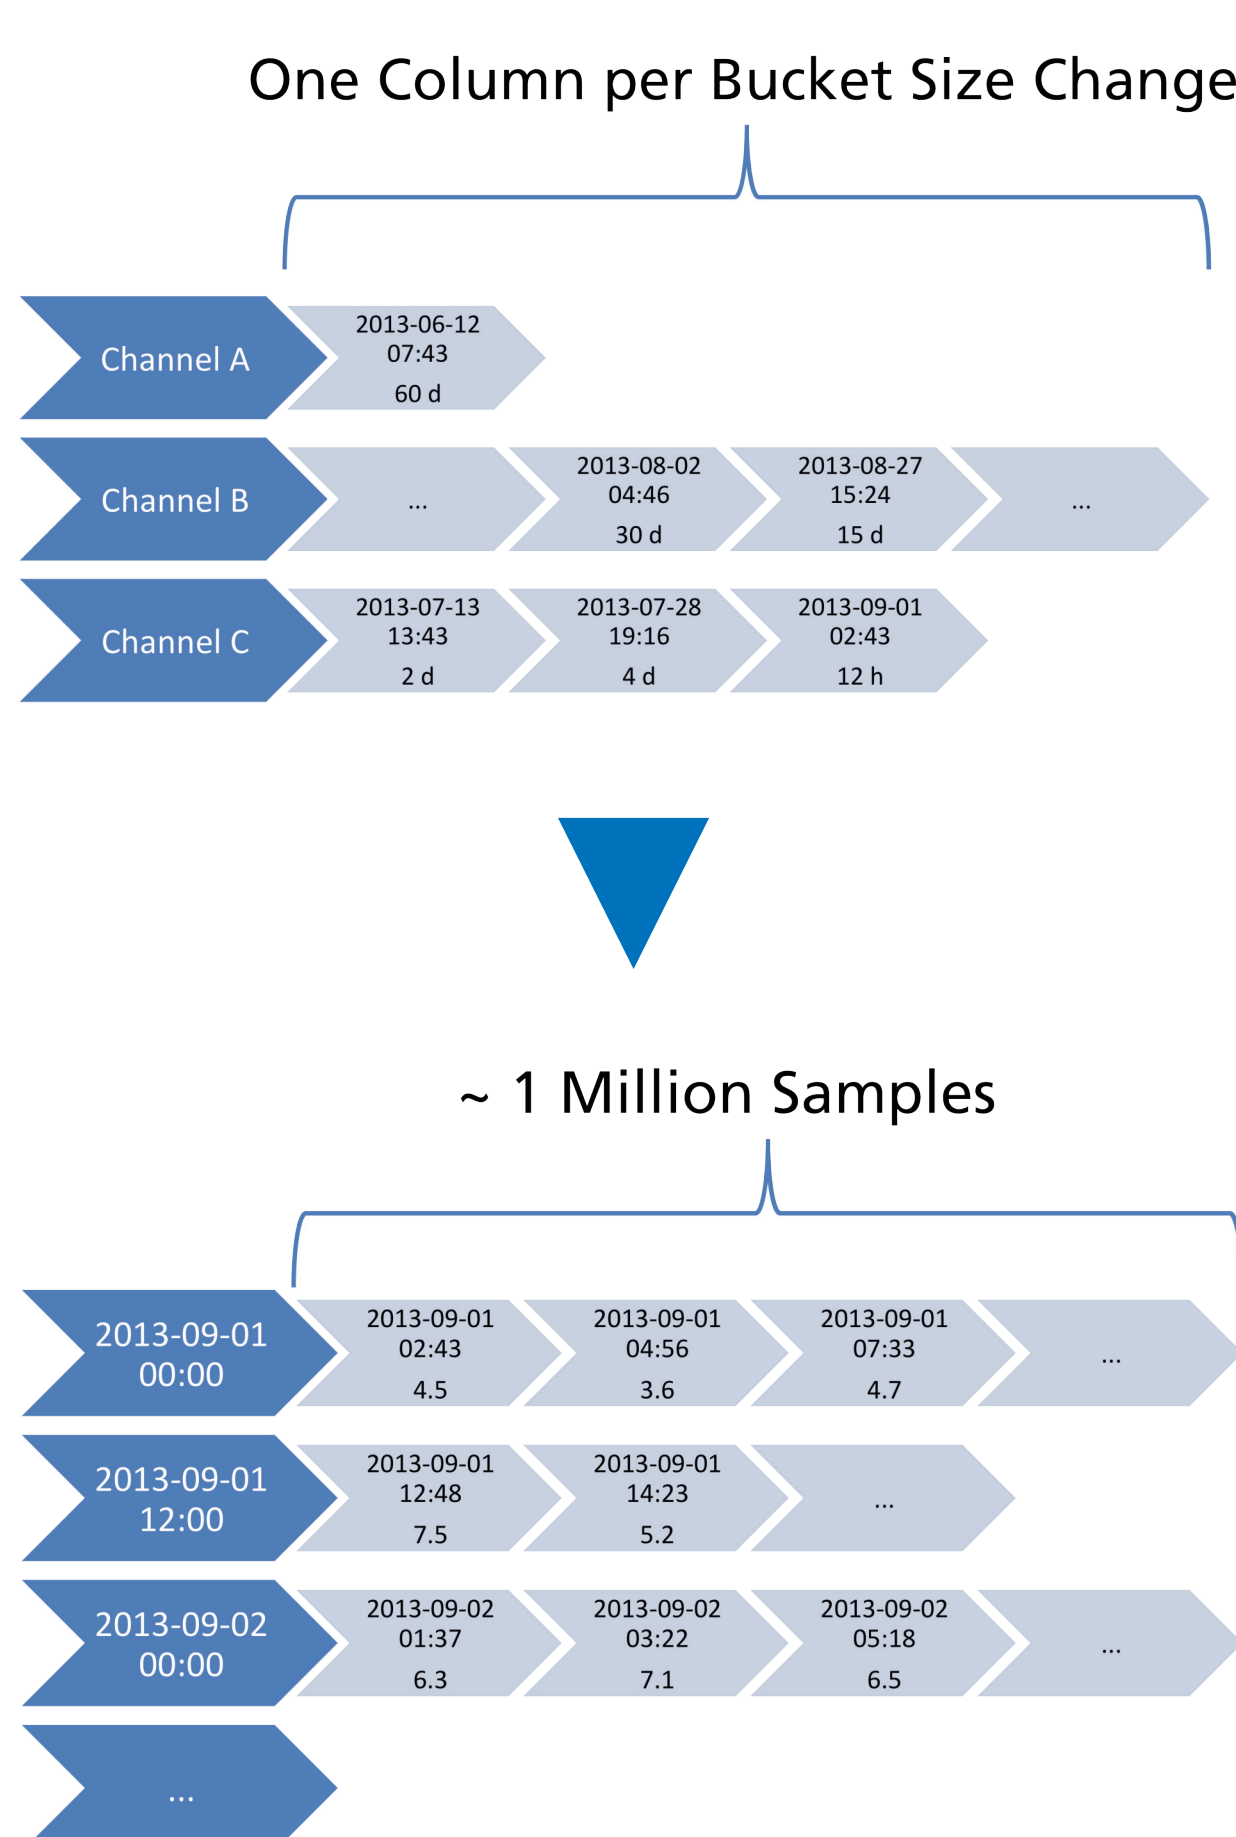
### Write Performance Linear Scaling



### Comparison with RDB Archiver



### Data Model




Download for free at  
<http://oss.aquenos.com/>

### Decimation




### Browsing


#### 3 Days




#### 1 Day




#### 4 Hours




#### 30 Minutes



#### 5 Minutes



#### 10 Seconds



### Apache Cassandra

#### History

- Initiated at Facebook
- Became an Apache Incubator project in 2009
- Apache top-level project since 2010

#### Data Model

- Similar to Google Bigtable
- Data is organized in column families
- Two-dimensional map that is stored in distributed hash-tables
- Data is replicated in the cluster

#### Benefits

- High I/O performance (sequential writes)
- Linear scaling
- Works with very large clusters (1000+ nodes)

