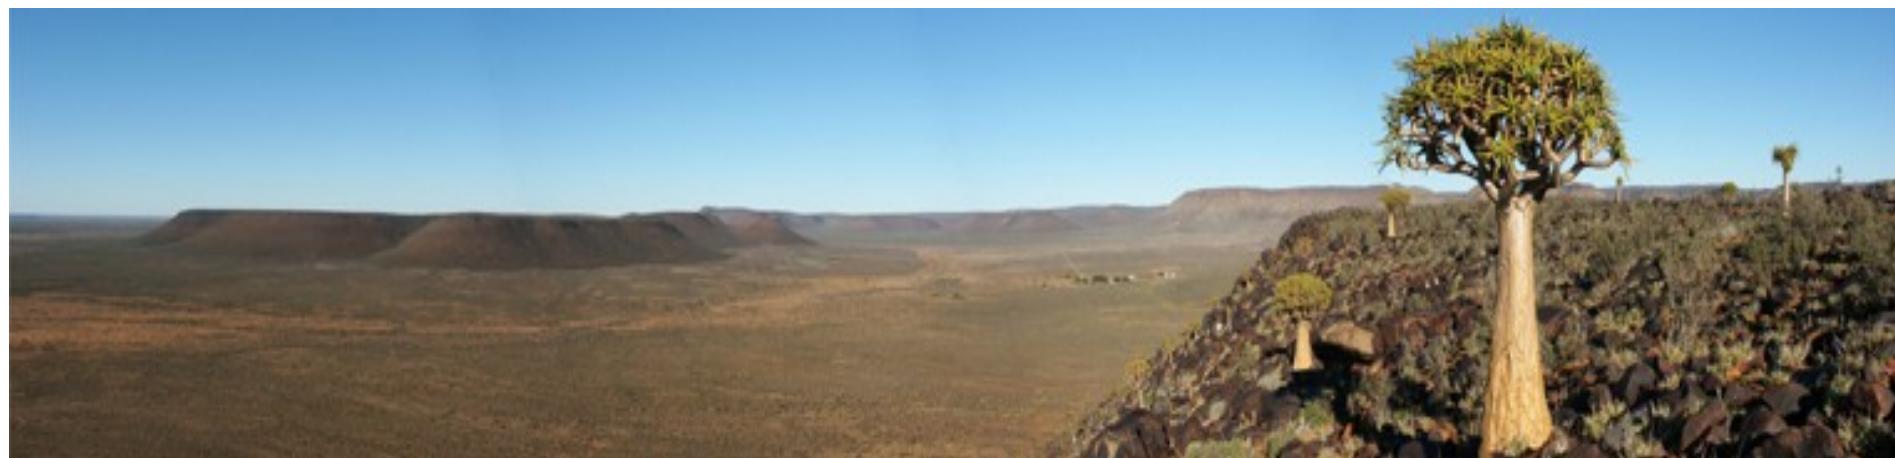


# *Virtualization and deployment management for the KAT-7 / MeerKAT control and monitoring system*

Neilen Marais

Senior Software Engineer

nmarais@ska.ac.za



# Introduction



- ◆ Radio Telescopes in the Karoo
  - ❖ KAT-7
  - ❖ MeerKAT
- ◆ Deployment of CAM subsystem
  - ❖ virtualization
  - ❖ automated deployment
  - ❖ share some experiences
- ◆ Work to improve started end 2011
  - ❖ Deployment fraught
  - ❖ Hardware failure -> extended downtimes
  - ❖ Limited development environments
    - quite different from deployments

# Requirements



- ◆ Deterministic+repeatable system configuration
- ◆ Versioned configuration history
  - ❖ quick revision roll-back/forward
- ◆ Minimize manual steps in deployment
- ◆ Minimize downtime
  - ❖ CAM software deployment
  - ❖ CAM system hardware failure
- ◆ Isolate resource usage on a shared server
- ◆ Easily deploy development environments
  - ❖ similar to site deployment environments
  - ❖ limited development hardware resource

# Shape of KAT-7 CAM

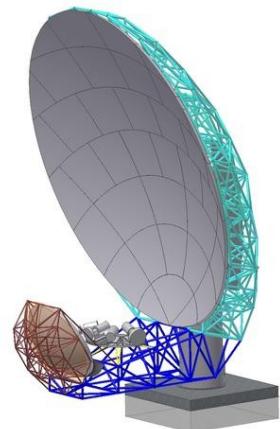
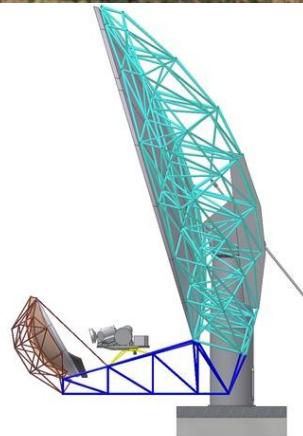
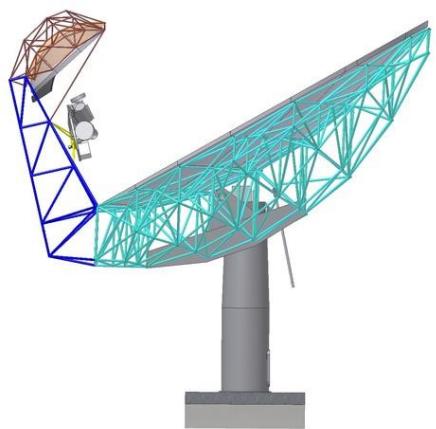


- ◆ Instrument is distributed
  - ❖ Karoo Array Telescope Control Protocol (KATCP)
  - ❖ Ethernet as fieldbus
- ◆ Telescope is Remote
  - ❖ Avoid human generated RFI
  - ❖ Control via high speed SANReN fibre (ring) network
- ◆ Operational Centre in Cape Town
  - ❖ Control room 700 km from site
  - ❖ Backup and long-term archiving
  - ❖ Development
- ◆ Mostly coded in the Python language

# KAT-7 Array



# Future MeerKAT Array



# KAT-7 Receptor Receiver



- ◆ Receiver Horn Antenna
- ◆ RF Low Noise Amplifier (LNA)
- ◆ Stirling Cryo cooler with Ion Pump
- ◆ RF Noise diode coupler
- ◆ RF Amplifier



# KAT-7 Receptor Pedestal



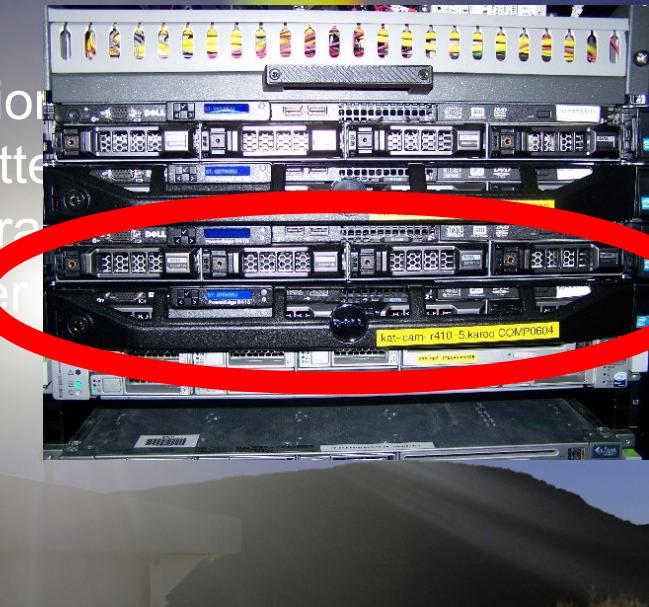
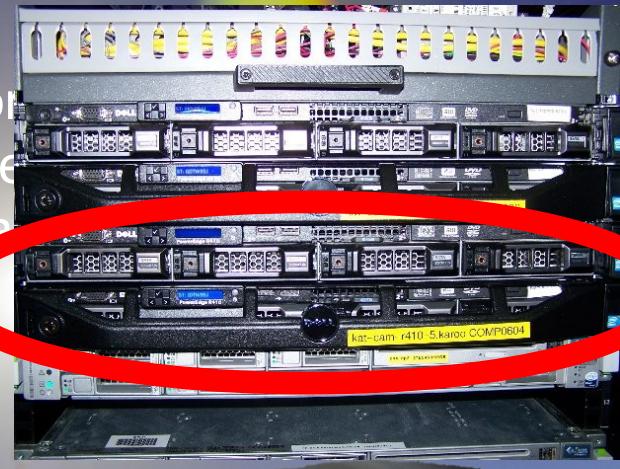
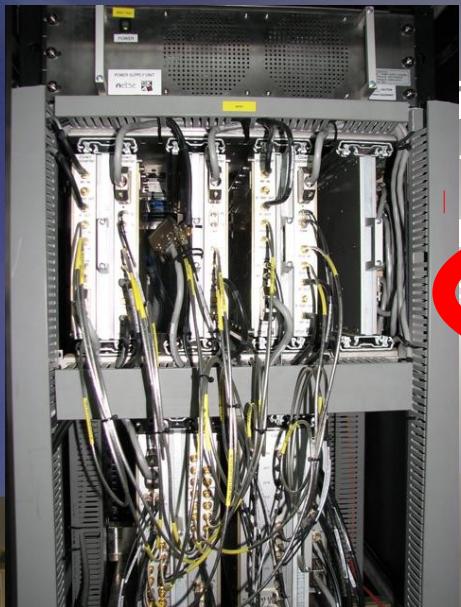
- ◆ Antenna positioner control unit
- ◆ RF amplifier/attenuator
- ◆ RF to optical transducer
- ◆ Pedestal chiller
- ◆ Building Management
- ◆ Weather station



# KAT-7 Compute Container



# Inside KAT-7 Compute Container



- ◆ Optical to RF transducers
- ◆ RF Down-conversion and conditioning
- ◆ FPGA based Digital Back End
- ◆ Data capture server
- ◆ Time/Frequency reference
- ◆ CAM Servers
- ◆ BMS
- ◆ Chiller

- Control Room



# Pieces of the Puzzle



- ◆ Server Virtualization
- ◆ Automated Deployment
- ◆ Deployment Configuration Database
- ◆ Combined: Hassle-free, deterministic, reliable deployment

# Virtualization



- ◆ Many Technologies, Many Makers
  - ❖ Full virtualization more flexible
  - ❖ Containerization more efficient
- ◆ Blurring of lines
- ◆ Other considerations
  - ❖ Familiarity
  - ❖ Licensing
  - ❖ Supported Host environments

# Server Virtualization: Proxmox VE



- ◆ Specialized Hypervisor distribution based on Debian GNU/Linux
  - ❖ FOSS licensing: no cost, no hassle
- ◆ Supports both:
  - ❖ Containers: OpenVZ
  - ❖ Full virtualization: KVM
- ◆ Simple and quick host install
- ◆ Easy to use web-based management tools
- ◆ Pre-configured base OS containers

# Performance



- ◆ CAM uses soft-realtime design
  - ❖ Only needs enough aggregate CPU throughput
- ◆ Similar aggregate CPU utilization on host before and after virtualization
- ◆ IO-bound tests using 10 GbE interface
- ◆ Using different virtualization options

## Test Machine

SUN FIRE X4150

2x Intel(R) Quad-core Xeon(R) E5450 CPUs

16 GB RAM

Gen 1 Myricom Myri10GE 10GbE

# Performance results



Config	Rate (Gb/s)	CPU use (%)	Relative rate (%)	CPU / Gb/s (%)
Baseline	5.49	65.8	100.0	12.0
Host	4.80	59.2	87.5	12.3
OVZ exclusive	4.65	61.2	84.7	13.2
OVZ veth	3.72	21.1	67.8	5.8
OVZ venet	3.86	20.3	70.4	5.3
KVM virtio	2.39	60.5	43.6	25.3

- ◆ Baseline: Ubuntu 10.04
- ◆ CAM uses veth
  - ❖ Most flexible
- ◆ Only 2x1Gb interfaces in production

# Desktop virtualization: Virtualbox



- ◆ Simulated system on developer, commissioner workstations
  - ❖ Toy-KAT VM
- ◆ Variety of workstation OSes
  - ❖ Can't take over whole machine for hypervisor
- ◆ Virtualbox virtualization host runs on them all
- ◆ Also FOSS licensing
- ◆ Not production use
  - ❖ positive experience

# Software Configuration Management



- ◆ Automated Deployment Scripts
- ◆ Deployment Configuration Database
- ◆ Configuration Manager

# SW Configuration Management Tech



- ◆ Preferred a Python based solution
- ◆ Considered existing 'full stack' systems (e.g. Puppet, Chef, Saltstack)
  - ❖ Central management server
  - ❖ Upfront time investment

# Our Sw Conf Management System



- ◆ Experimented with Fabric
  - ❖ Python based SSH automation library
- ◆ Script logic is defined in Python
  - ❖ Only requires SSH server on nodes to be managed
- ◆ Started capturing node configuration details
  - ❖ simple text file in INI-format
- ◆ Started implementing a Configuration Manager
  - ❖ Parsed the configuration database file
- ◆ Fabric library functions to deploy tasks to nodes
- ◆ Soon got team buy-in

# Configuration Database



- ◆ Node network configuration
- ◆ Node hosting type
- ◆ Assigned node resources
  - ❖ number of CPUs
  - ❖ RAM
  - ❖ Unique container / VM ID number
  - ❖ Diskspace, etc.
- ◆ Other meta information

# Disposability and Persistence



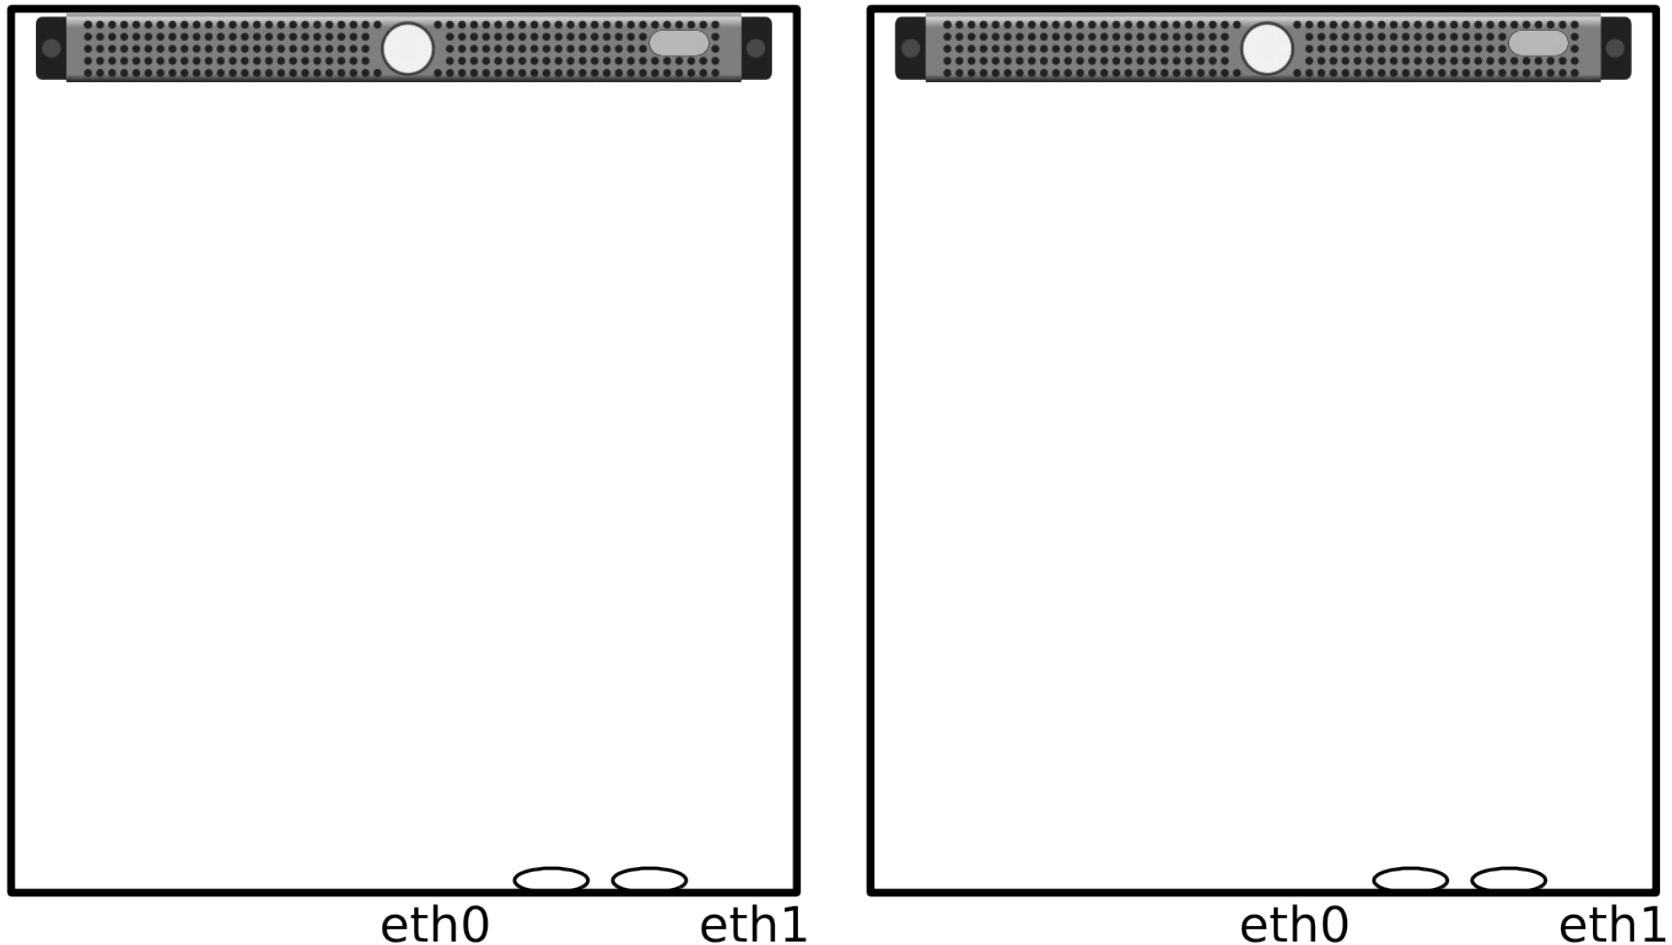
- ◆ Node containers are treated as disposable
  - ❖ Production containers rebuilt at each major release
- ◆ Development environments are routinely rebuilt
- ◆ Persistent data has to be managed separately
- ◆ NAS server, exported via NFS
  - ❖ Node NFS mounts configed as part of deployment
  - ❖ Potential issues with changing dB schemas
  - ❖ Central point for backups

# Deployment steps



- ◆ Configuring Proxmox hypervisor hosts
- ◆ Provisioning
- ◆ Configuration

# Unconfigured Hosts



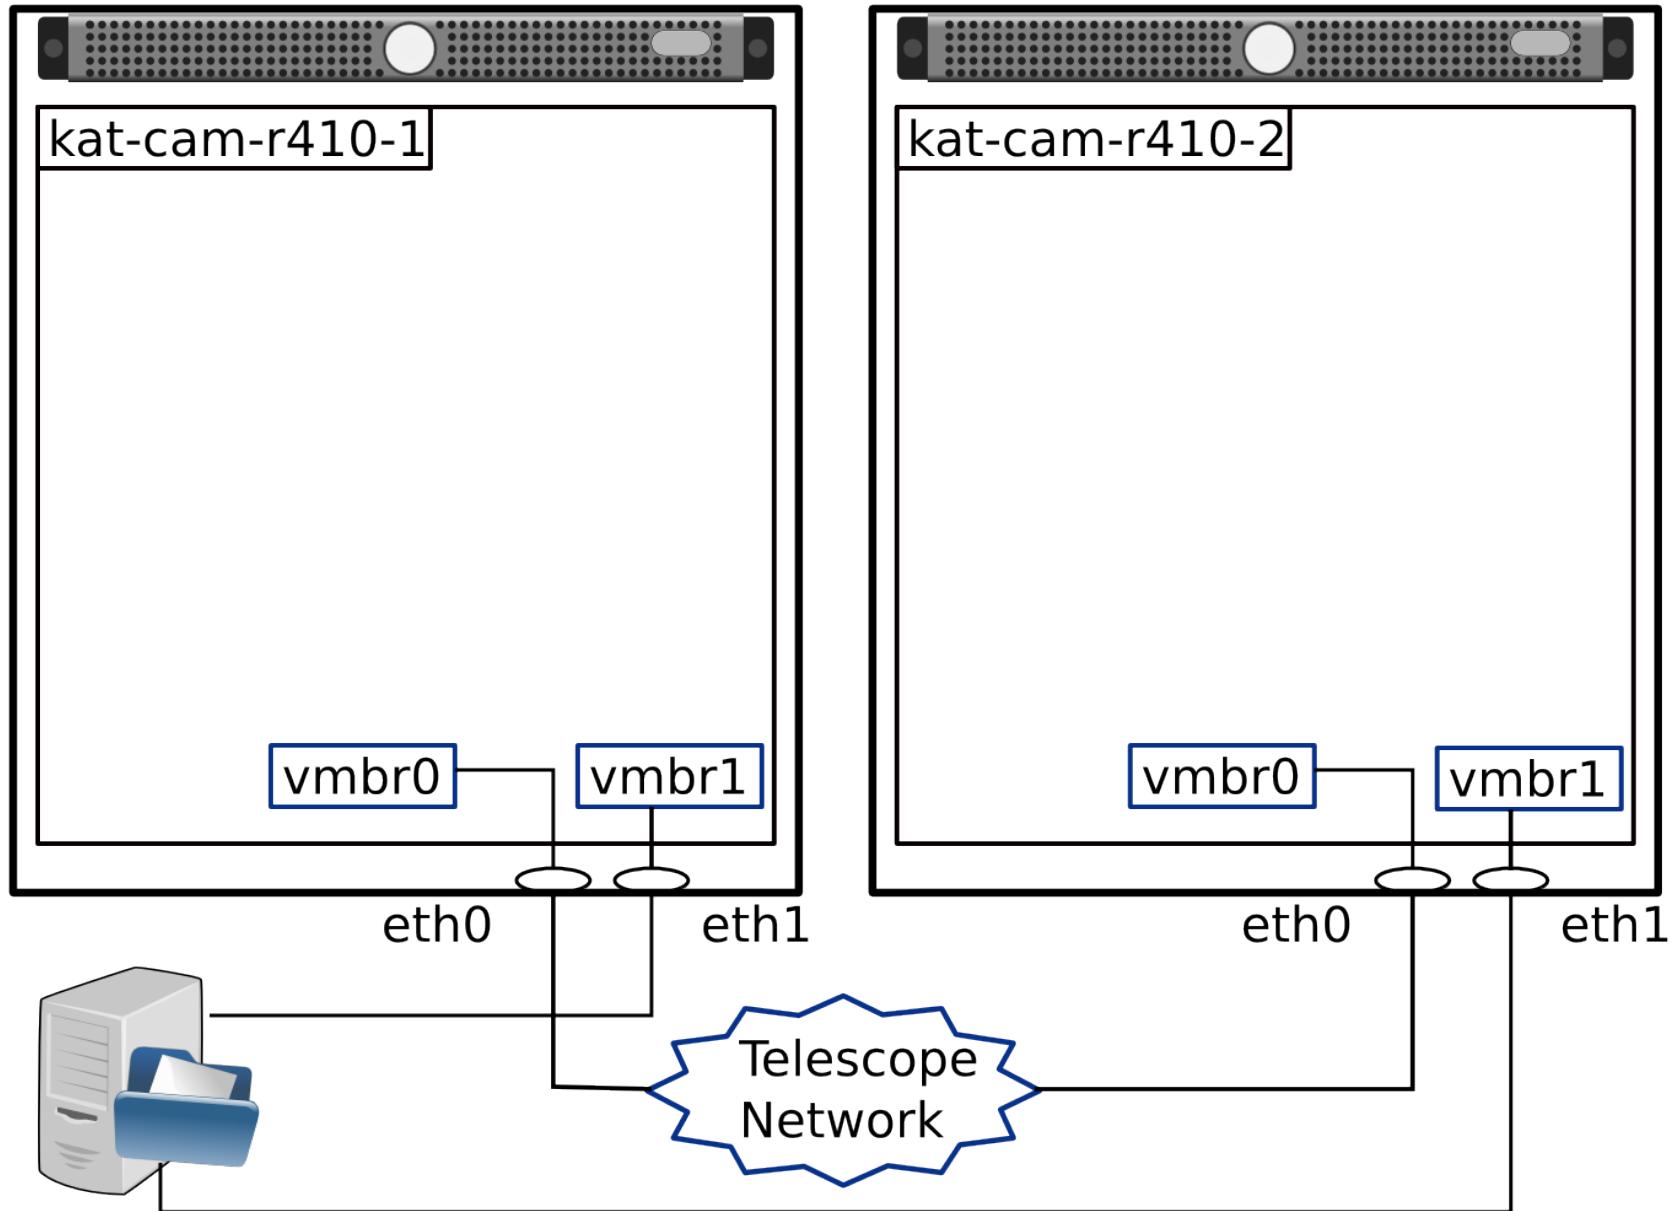
# Configuring Proxmox host



- ◆ Rarely done
- ◆ Install base Proxmox from CD
- ◆ Takes about 15 minutes including fab below

```
fab -H root@kat-cam-r410-1,root@kat-cam-r410-2\  
    proxmox.configure_host
```

# Proxmox installed and configured

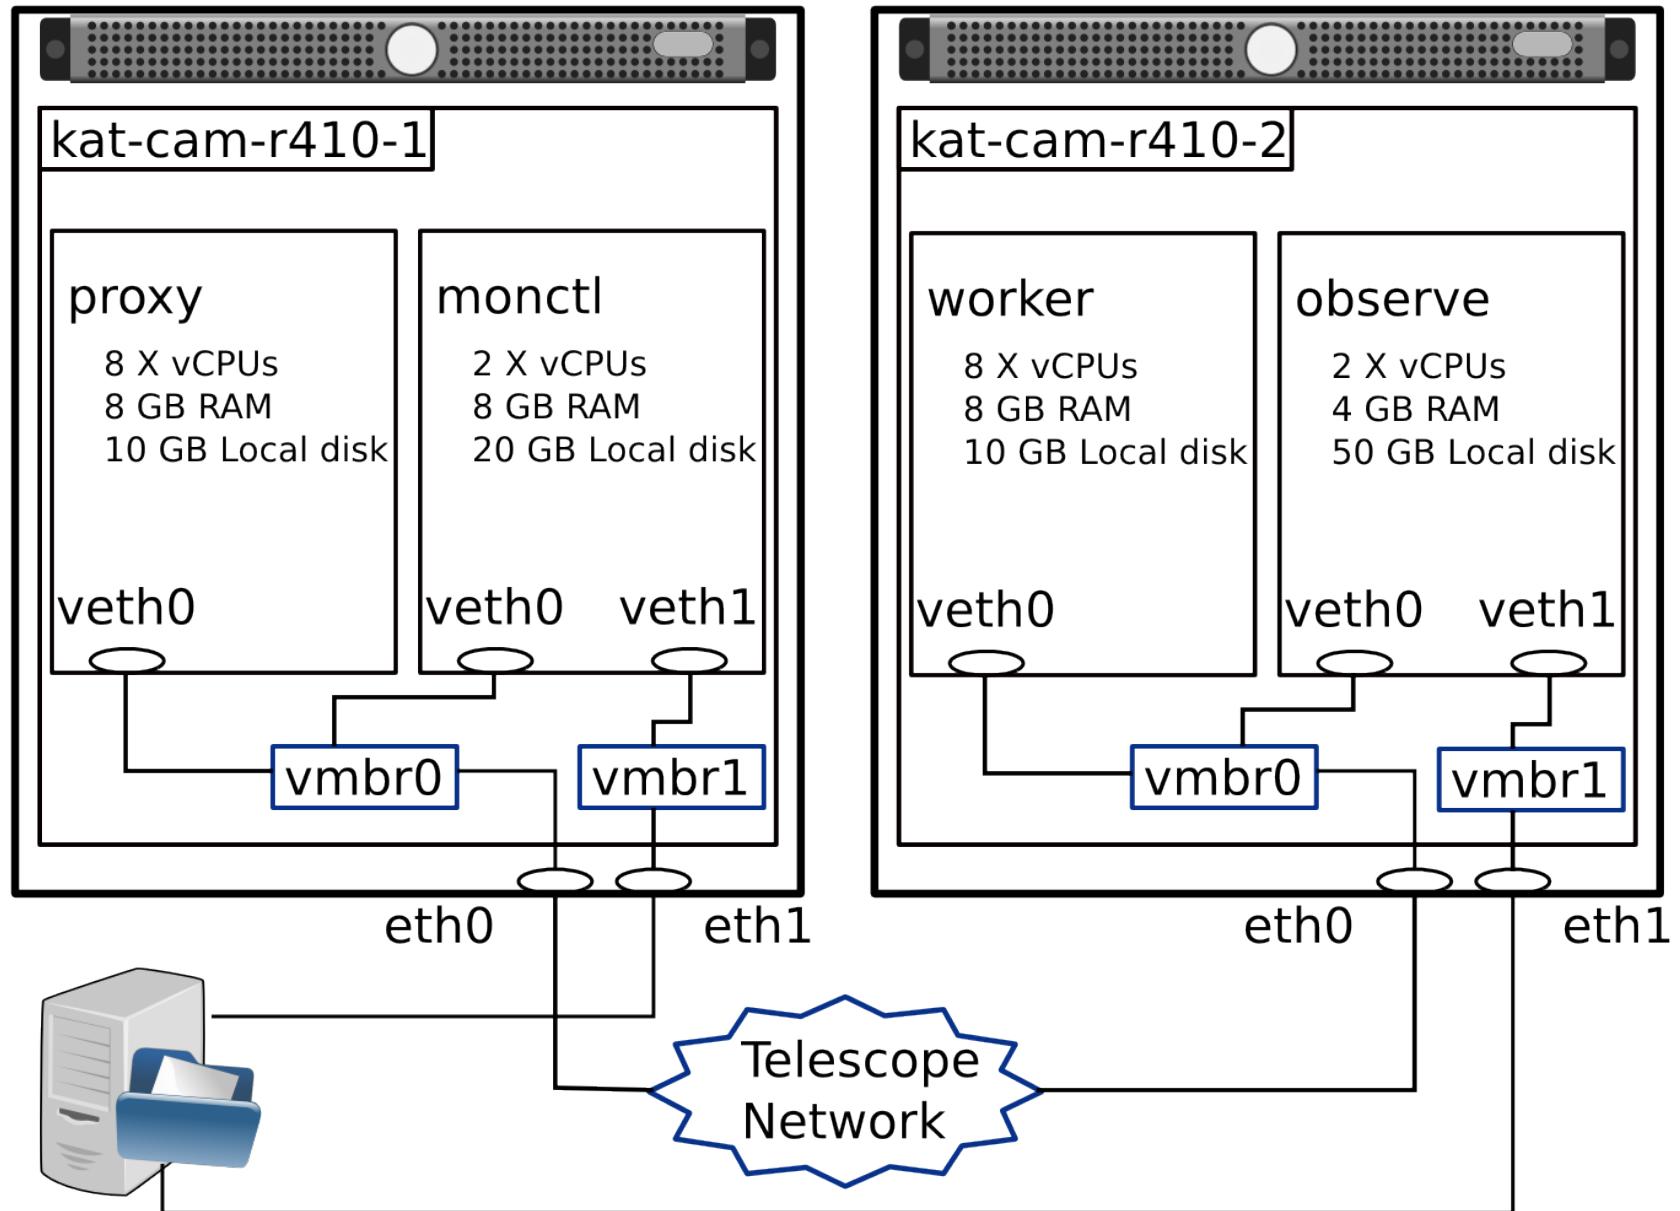


# Provision Node Containers



```
fab proxmox.create_containers_by_group:\  
karoo_system_nodes,700
```

# Virtual Node Containers Provisioned



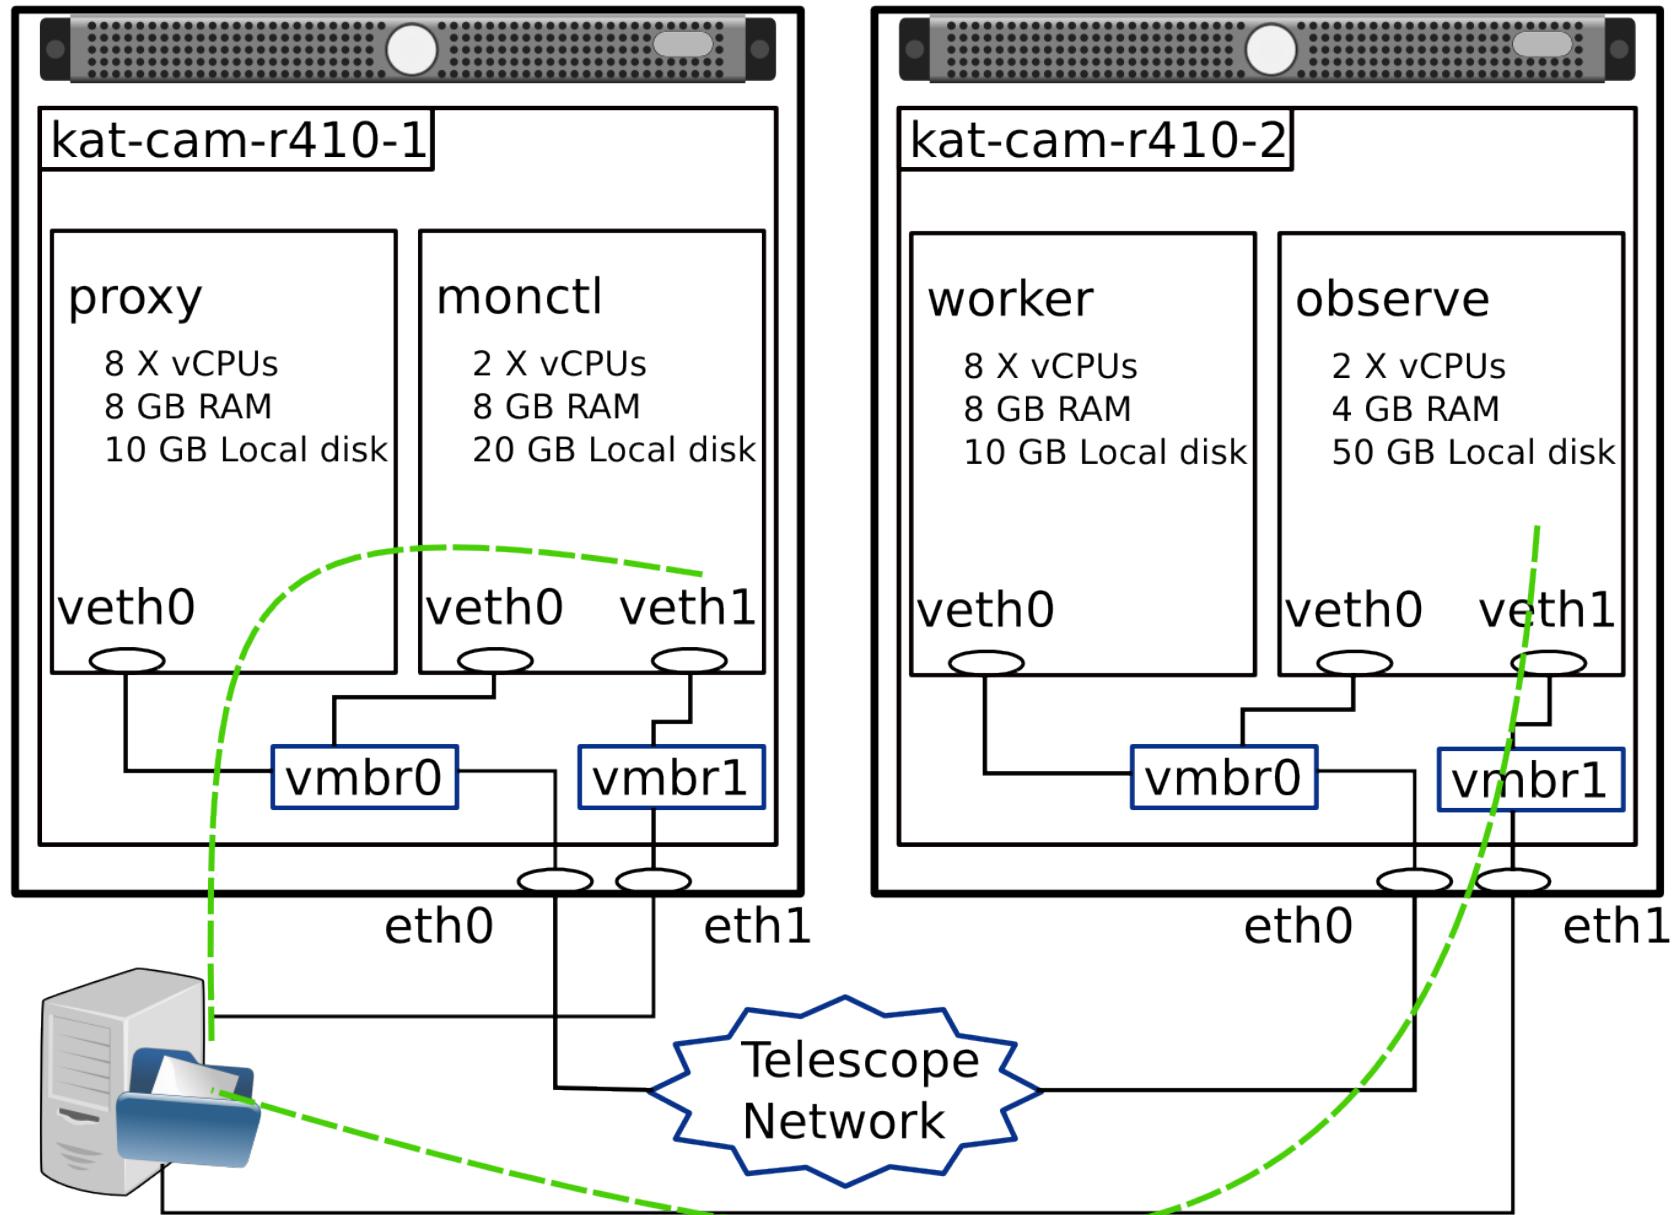
# Configure Nodes



```
fab kat_deploy.install_nodes_by_group:\  
karoocamv7-requirements.txt
```

- ◆ Install system packages (apt-get install ...)
- ◆ Install python packages (pip install ...)
- ◆ Configure NFS mounts
- ◆ Check out, build, install and configure packages from internal SVN
- ◆ Configure web servers, cron jobs, other OS level services as required
- ◆ Set up databases (schemas if needed)

# Nodes Configured



# Some Experiences



- ◆ Importance of transitioning from a mostly manual to automated deployment step by step.
- ◆ Remaining deployment problems are entered into our issue tracker
  - ❖ Deployment issues are prioritised for fixing
- ◆ Important to make deployment processes idempotent
- ◆ Important to make each step reliable
  - ❖ Local copies of internet based resources (PyPI and Ubuntu repositories)
  - ❖ Unexpected race conditions when things are not done at 'human' speed
- ◆ Usefulness of virtualization to allow testing and experiment with the deployment process -- you can just throw away and re-build a virtual node to test from-scrach deployment.
  - ❖ Also means we can test deployment, and not just our software

# Conclusion



Most frequently experienced advantages are:

- ◆ Easy deployment of realistic development/testing environments
  - ❖ including virtual networking mirroring actual configurations)
- ◆ Ability to quickly switch between software versions by switching containers
- ◆ Have largely met our goals

Future work:

- ◆ Deployment to fresh containers in the Continuous Integration process
  - ❖ Running full integration test suite on these containers
- ◆ Automatic daily building of Toy-KAT VMs
- ◆ Converting all legacy configuration scripts to the Fabric framework
- ◆ Future MeerKAT deployment should be more of the same
  - ❖ more complex network configuration

# Thank you!



nmarais@ska.ac.za

<http://www.ska.ac.za>